

Synchronous Dataflow

A Crash Course in Infinite Programs


Nic Hollingum

USYD

7 Apr, 2011

Problem

- ▶ Not all programs are suited to traditional Von-Neuman model
- ▶ Digital Signal Processing
- ▶ Programming these devices is difficult[Lee 87]
- ▶ Represent this kind of computation more naturally




- ▶ Termination is ill-defined, programs can run for infinity
- ▶ Expressive notion of data rates
- ▶ Control via sequencing of data - ties inputs to outputs
- ▶ hence Synchronous Dataflow

Tokens

- ▶ Single unit of data
- ▶ Data Streams represented by sequences of tokens
- ▶ Discreet, but vary in size
- ▶ Some tokens are provided by user at execution, others arrive as input
- ▶ Data-rates of tokens must be known at compile time



- ▶ individual computation units
 - ▶ Simple: $2 \rightarrow \text{add} \rightarrow 1$
 - ▶ Complicated: $2^n \rightarrow \text{FFT} \rightarrow 2^n$
- ▶ Consumes a number of tokens on input channels and produces on others
- ▶ amount consumed need not equal amount produced by predecessor
 - ▶ Must be known at compile time
 - ▶ Possibly consumes and produces from the same channel (feedback)
- ▶ Stateful or Stateless




Nic Hollingum



- ▶ Graphical representation of SDF network
- ▶ Actors are vertices
- ▶ Channels are Arcs
- ▶ Arcs are annotated with Delays and Consumption/production

SDF Graphs



[Batt 96]

Topology Matrix

- ▶ Graph: $\Gamma = \mathbb{R}^{|L| \times |A|}$
- ▶ A row for each link and a column for each actor

$$\Gamma_{l,a} = \begin{cases} 0 & \text{if } a \text{ is not an endpoint of } l \\ \text{prod}(l) & \text{if } a \text{ is the start-point of } l \\ -\text{cons}(l) & \text{if } a \text{ is the end-point of } l \end{cases}$$

Topology Matrix

$$\Gamma = \begin{bmatrix} 1 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & -1 & 0 & 0 & 0 & 0 \\ 0 & -2 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & -1 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & -2 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & -2 & 0 \\ 0 & 0 & 0 & 3 & 0 & -2 & 0 \\ 0 & 0 & 0 & 4 & 0 & 0 & -3 \\ 0 & 0 & 0 & 0 & 0 & 8 & -9 \end{bmatrix}$$

- $$\begin{bmatrix} 6 \\ 6 \\ 12 \\ 6 \\ 18 \\ 9 \\ 8 \end{bmatrix}$$

Pediodic Schedule

A list of actor invocations that return to the steady state

Pediodic Schedule

A list of actor invocations that return to the steady state

Init Schedule

A list of actor invocations that flood the buffers to the desired level

Pediodic Schedule

A list of actor invocations that return to the steady state

Init Schedule

A list of actor invocations that flood the buffers to the desired level

“Death Schedule”

A list of actor invocations that drain the buffers to the initial level (i.e. with the given delays)

Init and Death

SDF Crash Course

Nic Hollingum

Scheduling


- ▶ Fills and clears the buffers
- ▶ We have to know in advance how much we want to fill the buffers
- ▶ It may be a good idea to fill so that **any** periodic schedule may execute
- ▶ presume all actors simultaneously execute all their repetitions at the start
- ▶ delay actor invocations if we know buffers won't be full.
 - ▶ Strategically place actors
 - ▶ avoid costly dynamic scheduling



- ▶ massive virtual machines/multiple virtual machines
- ▶ network bandwidth is an important consideration
- ▶ small probability of node failure, whole computation is gone
- ▶ develop schemes to ensure fault tolerance
 - ▶ Statically replicate multiple nodes
 - ▶ Dynamically re-compute if a node fails
- ▶ examine trade-off between Makespan, Network Cost, and Fault tolerance

- ▶ abstraction from (single-core/multi-core/multi-processor/cloud) machines
- ▶ Use cost matrices to account for varying bandwidth/processing capabilities
 - ▶ 2 cores in the same machine have 0 communication cost
 - ▶ Actors assigned to a 1ghz processor take longer than those on a 1.8ghz processor
- ▶ Use this machine specification to generate a mapping of actors to processors
- ▶ Most must be approximated, optimality is NP-Hard

 S. S. Battacharyya, P. K. Murthy, and E. A. Lee.
Software Synthesis from Dataflow Graphs.
Kluwer Academic Publishers, 101 Philip Drive
Assinippi Park
Norwell, Massachusetts 02061 USA, first Ed., 1996.

 E. A. Lee and D. G. Messerschmitt.
“Static scheduling of synchronous data flow programs
for digital signal processing”.
Computers, IEEE Transactions on, Vol. 100, No. 1,
pp. 24–35, 1987.